



# Operating Systems – spring 2024

## Tutorial-Assignment 4

Instructor: Hans P. Reiser

<b>Submission Deadline: Monday, Februar 6th, 2023 – 23:59</b>
---

A new assignment will be published every week. It must be completed before its submission deadline (late policy for programming assignments: up to two days, 10% penalty/day)

**Lab Exercisess** are theory and programming exercises discussed in the lab class. They are not graded, but should help you solve the graded questions and prepare for the final exam. Make sure to read and think about possible solutions before the lab class.

**T-Questions** are theory homework assignments and need to be answered directly on Canvas (quiz).

**P-Questions** are programming assignments. Download the provided template from Canvas. Do not fiddle with the compiler flags. Submission instructions can be found in the introductory section below.

Topics of this assignment are threads and segmentation, and you will extend the scheduler from the previous assignment with priority-based scheduling.

Plan for this lab class

- Part 4.1 is revisiting the topic of segmentation from the lecture. It is not unlikely to find questions like the part (b) in the exam.
- Part 4.2 aims at practicing programming with pthreads. This is not yet needed for this week's programming assignment, but you will need this for a later assignment.
- In the second half of the lab class TA's are available to give advice on the programming assignment 4. Ideally you read the assignment and download the code templates to skel / to your PC before the class.

## Lab 4.1: Memory Management and Segmentation

- Explain the difference between a virtual and a physical address.
- How does segmentation work?
- Assume a system with 16-bit virtual addresses that supports four different segments, which uses the following segment table:

Segment Number	Base	Limit
0	0xdead	0x00ef
1	0xf154	0x013a
2	0x0000	0x0000
3	0x0000	0x3fff

Complete the following table and explain briefly how you derived your solution for each row in the table.

Virtual Address	Segment Number	Offset	Valid?	Physical Address
	3	0x3999		
0x2020				
		0x0204	yes	
			yes	0xf15f

## Lab 4.2: Programming with pthreads

- Write a small program that creates five threads using the pthread library.  
Each thread should count down in one second steps from 3 to 0, printing a message each second (e.g. This is thread 4, 3 seconds remaining...)  
See `man 3 sleep` for how to wait for one second.  
Note: This part is not (yet) relevant for this week's programming assignments, but we will come back to using pthreads later (in another programming assignment).

## T-Question 4.1: Threads and Thread Models

- a. Assume one thread of an application is currently executing on the (single) CPU. What events can cause the activation of a different thread of the same application?

There may be differences between User-Level Threads (ULT) and Kernel-Level Threads (KLT), so answer separately for those two.

4 T-pt

Event to consider	ULT	KLT)
The thread invokes a blocking system call (for example, reading from a file, which has to wait for the disk)		
The thread causes a "segmentation fault" exception (with default fault handling).		
A timer interrupt occurs		
A thread voluntarily relinquishes the processor with a yield() operation		

Yes / No / Depends – more detailed answers on Canvas quiz.

- b. Which of the following objects are generally all shared by the threads of a process (i.e., managed by/stored in the PCB, not the individual TCB)?

1 T-pt

true    not true

- ☐    ☐    Code and stack
- ☐    ☐    Instruction pointer, register contents, open files
- ☐    ☐    Code, heap, open files
- ☐    ☐    Code, stack, open network connections
- ☐    ☐    Thread control block

- c. Do you agree with the statement "Switching to a different process is usually much faster than switching to a different user level thread, because switching processes can very efficiently be handled by the CPU scheduler of the operating systems"? Justify!

2 T-pt

## T-Question 4.2: Segmentation

- a. Assume a system with 16-bit virtual addresses that supports four different segments (the two most significant bits of the virtual address encode the segment number, the remaining 14 bits the offset within the segment). Assume we use the following segment table:

4 T-pt

Segment Number	Base	Limit
0	0x1000	0x15f3
1	0x8000	0x00FF
2	0x25f3	0x1000
3	0x4500	0x0300

Complete the following table.

Virtual Address	Segment Number	Offset	Valid?	Physical Address
0xC0DE				
	1	0x0200		
			yes	0x25f4

## P-Question 4.1: Priority Scheduler

Download the template **p1** for this assignment from Canvas. You may only modify and upload the file `scheduler.c`.

Priority scheduling assigns each scheduling entity (i.e., a process or thread) a priority. For each priority the scheduler has a ready queue into which ready threads with the respective priority are enqueued. The scheduler always selects the first thread from the non-empty ready queue of the highest priority. If multiple threads have the same priority (i.e., a queue contains more than one thread), the scheduler employs round robin scheduling within the queue. Refer to the lecture slides for more details.

In this assignment, you will replace the simple FIFO scheduler from assignment 3 with a more complex priority scheduler. We assume in the following that all thread control blocks are stored in a static array. The thread ID is the index of a thread in that array. The scheduler queues store only the thread ID.

You can reuse the queue implementation from the previous assignment. A sample solution for that queue will be published after the late submission deadline (Wednesday 23:59).

- a. Implement/modify the event handler functions, which set the supplied thread's state and if necessary add the thread to the appropriate ready queue. **2 P-pt**

- `void onThreadReady(int threadId)` is called if a thread in waiting state becomes ready (e.g., thread was blocked on an I/O operation before, and the I/O operation has finished). The thread needs to be placed in the appropriate runqueue.
- `void onThreadPreempted(int threadId)` is called if a running thread was preempted. It also needs to be placed in the right runqueue, as it is ready to continue.
- `void onThreadWaiting(int threadId)` is called when a thread blocks (e.g., on an I/O operation). Such a thread enters the waiting state and will not be part of any runqueue.

- b. Implement the priority scheduling policy Your scheduling function should perform the following basic operations whenever a new thread needs to be selected: **4 P-pt**

- Find the ready queue with the highest priority that contains a ready thread
- Remove the first thread from the queue, updates its state and returns its thread id

- c. Add starvation prevention to the scheduler using the following additional rule: **4 P-pt**

- If a thread with priority  $P$  has been selected for four times without giving a lower-priority thread the chance to run, then instead of selecting a thread with priority  $P$  again, the scheduler will resort to the next lower priority queue that (1) is not empty and (2) does not break this starvation rule (i.e., exceeding the 4 times maximum without scheduling lower priority threads).
- Example: Assume that 0 is the highest, 10 is the lowest priority, and you have threads with priorities 7, 6 and 4; these threads do not block, i.e. after running they immediately return to the ready queue. In this case, the schedule will be:  
"4 4 4 4 6 4 4 4 4 6 4 4 4 4 6 4 4 4 4 7 4 ..."

*Hints:* The simplest solution to the starvation prevention will use a recursive approach to determine the right queue for thread selection.

```
int scheduleNextThread();
```